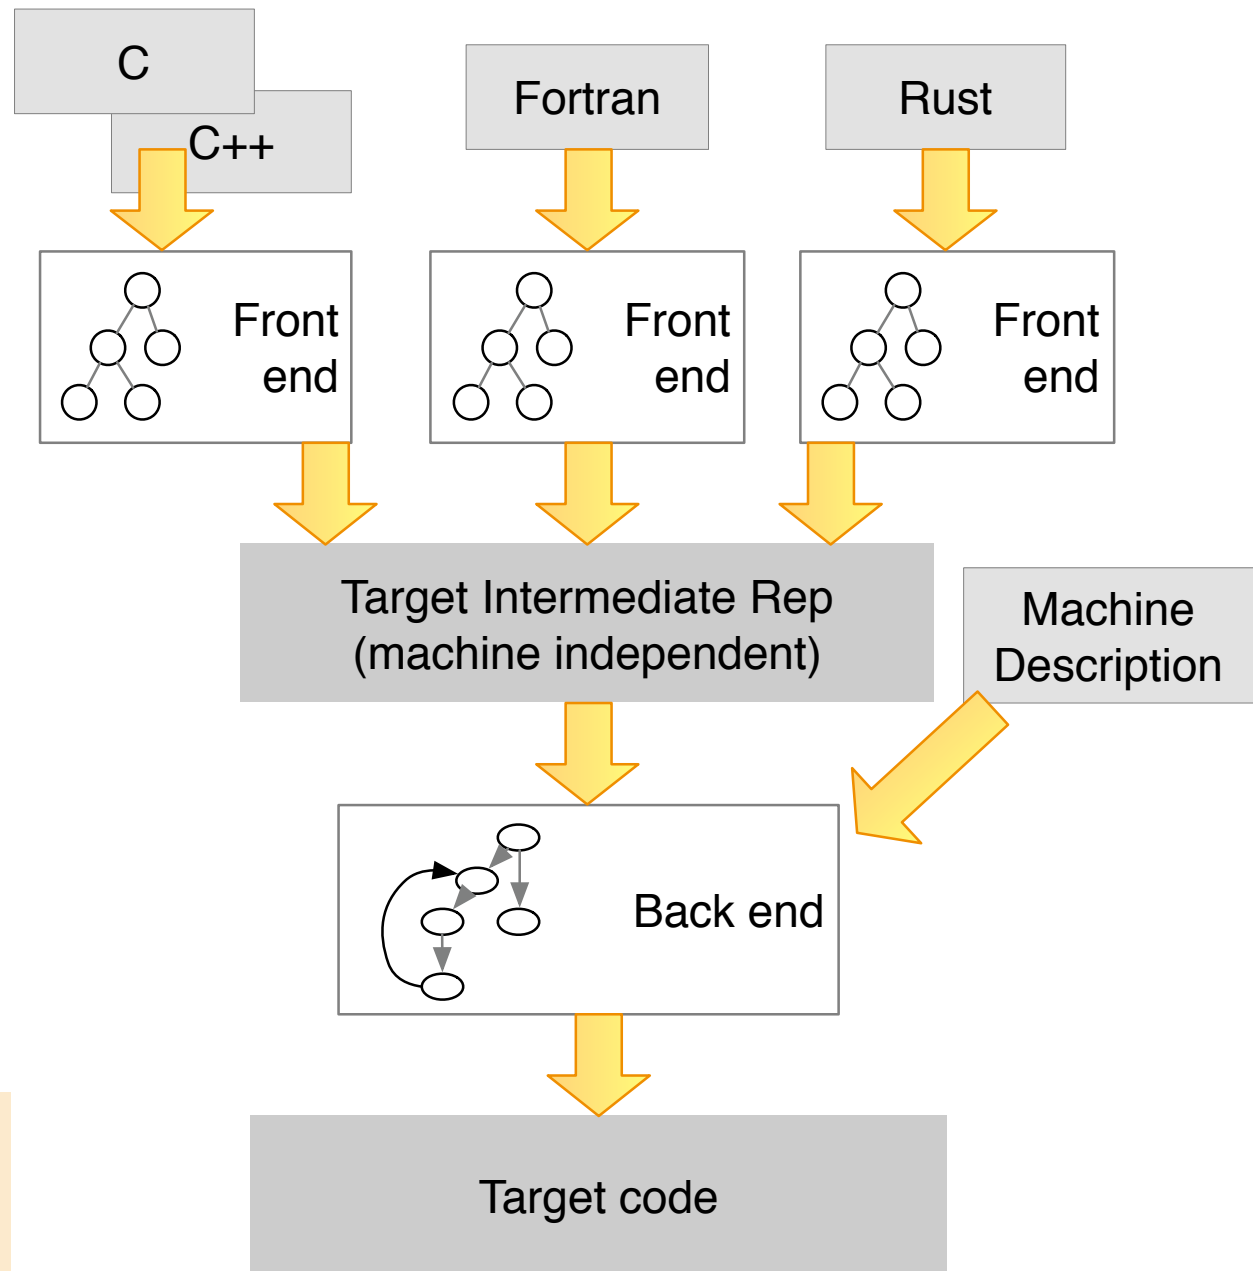


# *From Abstract Machine to Concrete Machine*

Generating Assembly Code  
(although we will generate C)



# Intermediate & target code



Example back ends:  
GCC, LLVM

“Sandboxed” languages  
look a little different ...  
VM instead of back end



# *Resources to Manage*

## Memory

Divided into text (code), stack, and heap

- Code: where the program code lives
- Stack: activation records; global data may also live here, or in its own area
- Heap: dynamically managed memory

## Registers

General-purpose and dedicated

Untyped (except for floating-point)

- But llvm registers (an infinite supply) are sorta kinda typed



# Registers and Memory

Memory (RAM) is large and slow

Much, much slower than modern processors

Registers are few but fast

They are *part of the processor*

As few as 4 general purpose registers (x86)

As many as 256 (Sparc)

16 to 32 is common

additional special purpose registers (program counter, flags register, etc.)

RISC machines are load/store architectures

No other operations on memory

Since we will generate C code, we will not do register allocation, but you should understand registers vs. memory.



# *Managing Registers and Memory*

Each architecture (processor + OS) has a set of *calling conventions*

How do I allocate a stack frame?

How do I pass parameters and results?

Which registers are preserved?

- Preserved: Callee must save, or leave alone
- Volatile: Callee may clobber; caller can save

Conventions described in “application binary interface” document



# Calling Convention - Skeleton

## Before the call:

- Put parameters into registers, or on the stack

- Save any volatile registers with important stuff

## On entry:

- Allocate a stack frame

- Save return address and old stack pointer

## On exit

- Place return value in register (or on stack)

- Restore stack frame

- Jump to return address

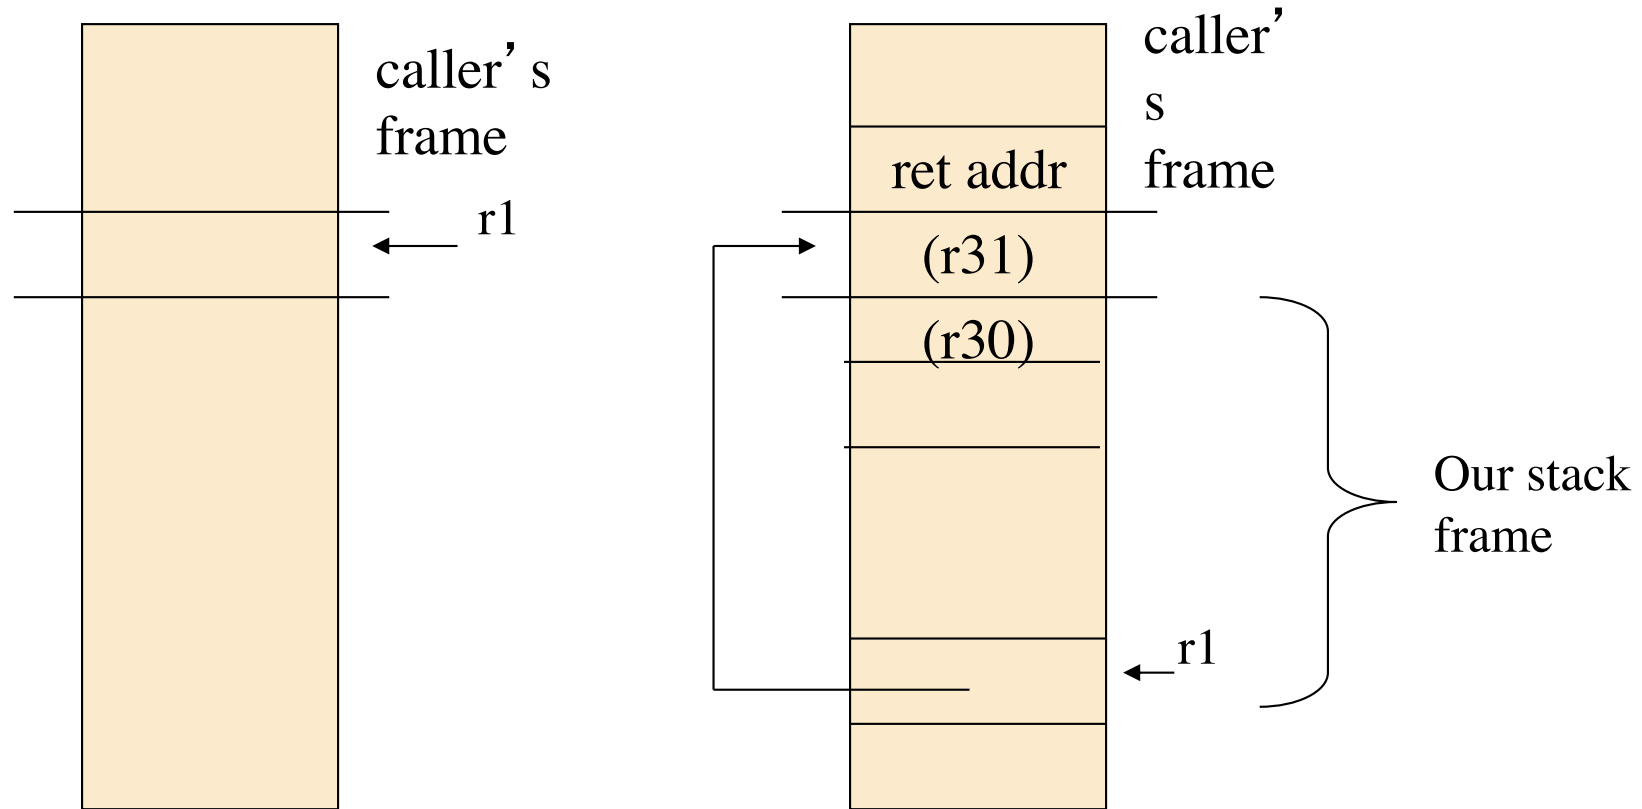
In assembly language, we might write macros (like inline functions, or C preprocessor macros) to handle the call/return conventions.



# Creating an Activation Record

*PowerPC example. There are probably mistakes in this.*

```
mflr r0          ; R0 <- return address
stmw r30,-8(r1)   ; Save R30 and R31 in my area
stw r0,8(r1)      ; Save return address into caller frame
stwu r1,-64(r1)   ; Allocate 64 bytes and save old sp
mr r30,r1         ; R30 <- R1 (huh? frame pointer?)
```



# Calling Conventions

Basic calling conventions are mostly language independent

Though you'll need a couple of implicit parameters, like "self-object"


Gcc will spill the beans with -S option

sample.c

```
static int main( ... ) { ... }  
  
void minimal( ) { return; }  
  
int sum(int a, int b) { .... }
```



gcc -S



sample.s

```
.section __TEXT,__ ...  
.align 2  
_sum:  
    stmw r30,-8(r1)  
    stwu r1,-64(r1)  
    mr r30,r1  
    stw r3,88(r30)  
    stw r4,92(r30)  
    lwz r0,88(r30)
```





# Calling Conventions: LLVM

Similar trick as gcc -S:

Use Clang or other llvm-targeted compiler to  
generate llvm intermediate code

Identify boilerplate and idioms

Generating C, we won't need this. We will translate a method call into a C function call.

If you generate assembly code, for this project or others, gcc -S is your friend. Write tiny functions with variations.



# *Tactic for Identifying Boilerplate*

## Two procedures

One does almost nothing

- and so should contain almost only boilerplate

One does a little something

- and so should have some code besides boilerplate

## Compare side-by-side

What is boilerplate on entry?

What is boilerplate on exit?



# PowerPC (old Mac OS X)

```
int sum(int a, int b) {  
    int c, d, e;  
    c = a;  
    d = b;  
    e = a + b;  
    return e;  
}
```

```
int nothing() {  
    return;  
}
```

```
section __TEXT,__text, ..., .align 2  
_sum:
```

```
    stmw r30,-8(r1)  
    stwu r1,-64(r1)  
    mr r30,r1
```

```
    stw r3,88(r30)
```

```
    stw r4,92(r30)
```

```
    ... code in sum but not in nothing ...
```

```
    lwz r0,40(r30)
```

```
    mr r3,r0
```

```
    lwz r1,0(r1)
```

```
    lmw r30,-8(r1)
```

```
    blr
```

```
    .align 2
```

```
    .globl _nothing
```

```
.section
```

```
    __TEXT,__text,regular,pure_instructions
```

```
    .align 2
```

```
_nothing:
```



# Sparc (old ix)

```
int sum(int a, int b) {  
    int c;  
    c = a + b;  
    return c;  
}
```

```
.section      ".text"  
.align 4  
.global sum  
.type  sum, #function  
.proc  04
```

sum:

save %sp, -120, %sp

st %i0, [%fp+68]

st %i1, [%fp+72]

ld [%fp+68], %i5

ld [%fp+72], %g1

add %i5, %g1, %g1

st %g1, [%fp-20]

ld [%fp-20], %g1

mov %g1, %i0

ret

restore

.size sum, .-sum

Shift register window  
(peculiar to sparc)

Note incredibly stupid  
code from gcc. Don't be  
too depressed by yours!

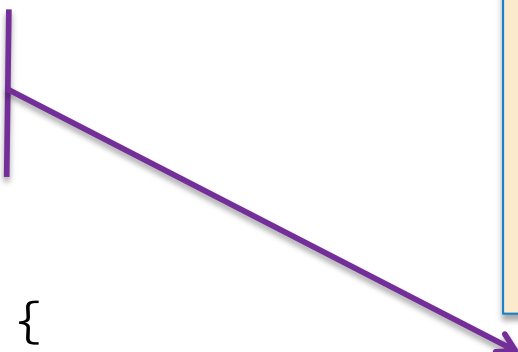
Branch delay slot  
(peculiar to sparc)



# clang (llvm) on MacOS X, x86:

```
int sum(int a, int b) {  
    int c;  
    c = a + b;  
    return c;  
}
```

```
int nothing() {  
    return 42;  
}
```



```
_sum:                                ## @sum  
    .cfi_startproc  
## BB#0:  
    pushq %rbp  
Ltmp2:  
    .cfi_def_cfa_offset 16  
Ltmp3:  
    .cfi_offset %rbp, -16  
    movq %rsp, %rbp  
Ltmp4:  
    .cfi_def_cfa_register %rbp  
    movl %edi, -4(%rbp)  
    movl %esi, -8(%rbp)  
    movl -4(%rbp), %esi  
    addl -8(%rbp), %esi  
    movl %esi, -12(%rbp)  
    movl -12(%rbp), %eax  
    popq %rbp  
    retq  
    .cfi_endproc
```



# Viewing llvm intermediate code

*clang -S -flto -emit-llvm sample.c  
more sample.s*

*When you see what clang or gcc emits, you  
won't feel so bad the code you generate.*

```
; Function Attrs: nounwind ssp uwtable
define i32 @sum(i32 %a, i32 %b) #0 {
    %1 = alloca i32, align 4
    %2 = alloca i32, align 4
    %c = alloca i32, align 4
    store i32 %a, i32* %1, align 4
    store i32 %b, i32* %2, align 4
    %3 = load i32* %1, align 4
    %4 = load i32* %2, align 4
    %5 = add nsw i32 %3, %4
    store i32 %5, i32* %c, align 4
    %6 = load i32* %c, align 4
    ret i32 %6
}
```

Allocate space in stack

Copy args into stack (why?)

Load arg values into  
registers

Add (result in register)

Save into 'c' in stack

Load 'c' into a register  
Return it



# Addressing Data

## Local variables

Offset from SP (stack pointer) or FP (frame pointer)

- Including variables declared in nested blocks. Do you see how?

## Object fields

Offset from object pointer

- Keep “this” object in register. Implicit parameter?
- Others: object reference is pointer

## Non-local

- “Display” or “lexical chain”
- (Not in Quack --- we don’t have nested block scopes!)

In C:

Objects are structs ... fields are members of struct

Classes are structs ... methods are function pointers

Local variables are C variables (with generated names)



# *For each method ...*

## Calculate total size of frame

### Save areas + links + local variables

- Easy in Quack: All local variables are 1 word (4 bytes); larger objects are pointers to heap

## Associate address with each local

### Offset within frame

- From stack pointer
- Or from frame pointer (fixed location in frame)

In C: We just create local variable in generated code.  
gcc or llvm will calculate frame sizes for us

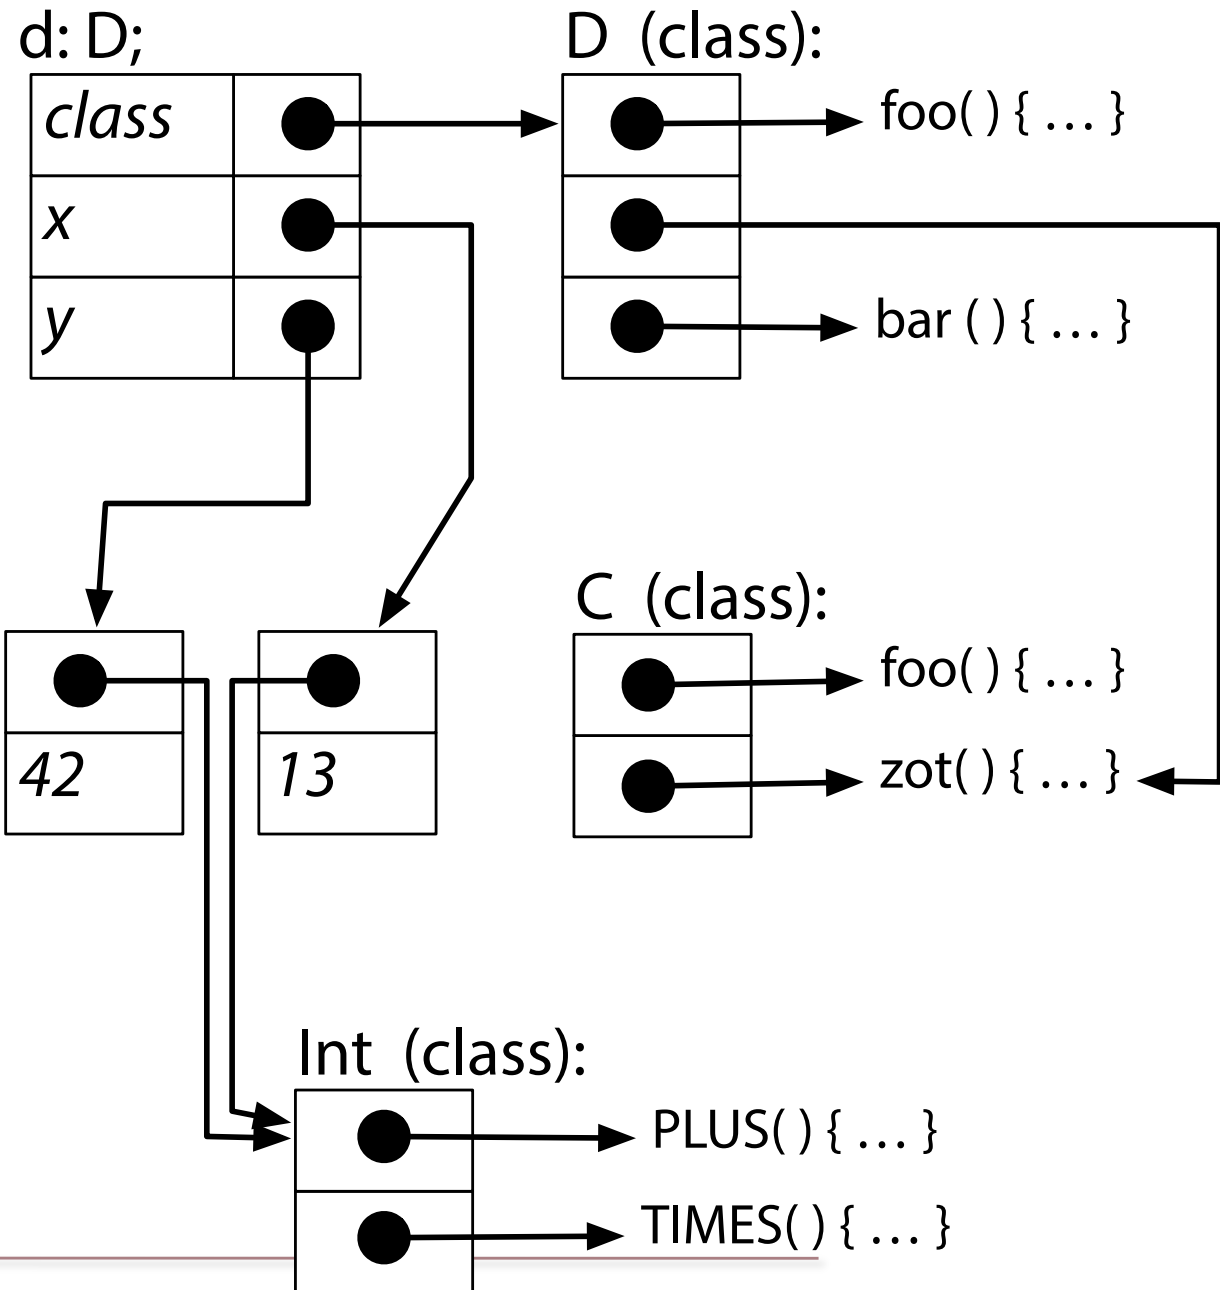




# Addressing fields and methods

```
class C( ) {  
    this.x = 42;  
    def foo( ): Int {  
        this.zot( );  
    }  
    def zot( ) { ... }  
}
```

```
class D() extends C {  
    this.x = 43;  
    this.y = 13;  
    def foo( ): Int {  
        this.zot( );  
        this.bar( );  
    }  
    def bar( ) { ... }  
}
```



# What does this really look like?

```
struct class_Pt_struct;
typedef struct class_Pt_struct* class_Pt;

typedef struct obj_Pt_struct {
    class_Pt clazz;
    obj_Int x;
    obj_Int y;
} * obj_Pt;

struct class_Pt_struct the_class_Pt_struct;

struct class_Pt_struct {
    obj_Pt (*constructor) (obj_Int, obj_Int );
    obj_String (*STRING) (obj_Obj);
    obj_Pt (*PRINT) (obj_Pt);
    obj_Boolean (*EQUALS) (obj_Obj, obj_Obj);
    obj_Pt (*PLUS) (obj_Pt, obj_Pt);
};
```


```
extern class_Pt the_class_Pt;

obj_Pt Pt_method_PLUS(obj_Pt this, obj_Pt other) {
    obj_Int this_x = this->x;
    obj_Int other_x = other->x;
    obj_Int this_y = this->y;
    obj_Int other_y = other->y;
    obj_Int x_sum = this_x->clazz->PLUS(this_x, other_x);
    obj_Int y_sum = this_y->clazz->PLUS(this_y, other_y);
    return the_class_Pt->constructor(x_sum, y_sum);
}

/* The Pt Class (a singleton) */
struct class_Pt_struct the_class_Pt_struct =
{
    new_Pt, /* Constructor */
    Obj_method_STRING,
    Pt_method_PRINT,
    Obj_method_EQUALS,
    Pt_method_PLUS
};
```

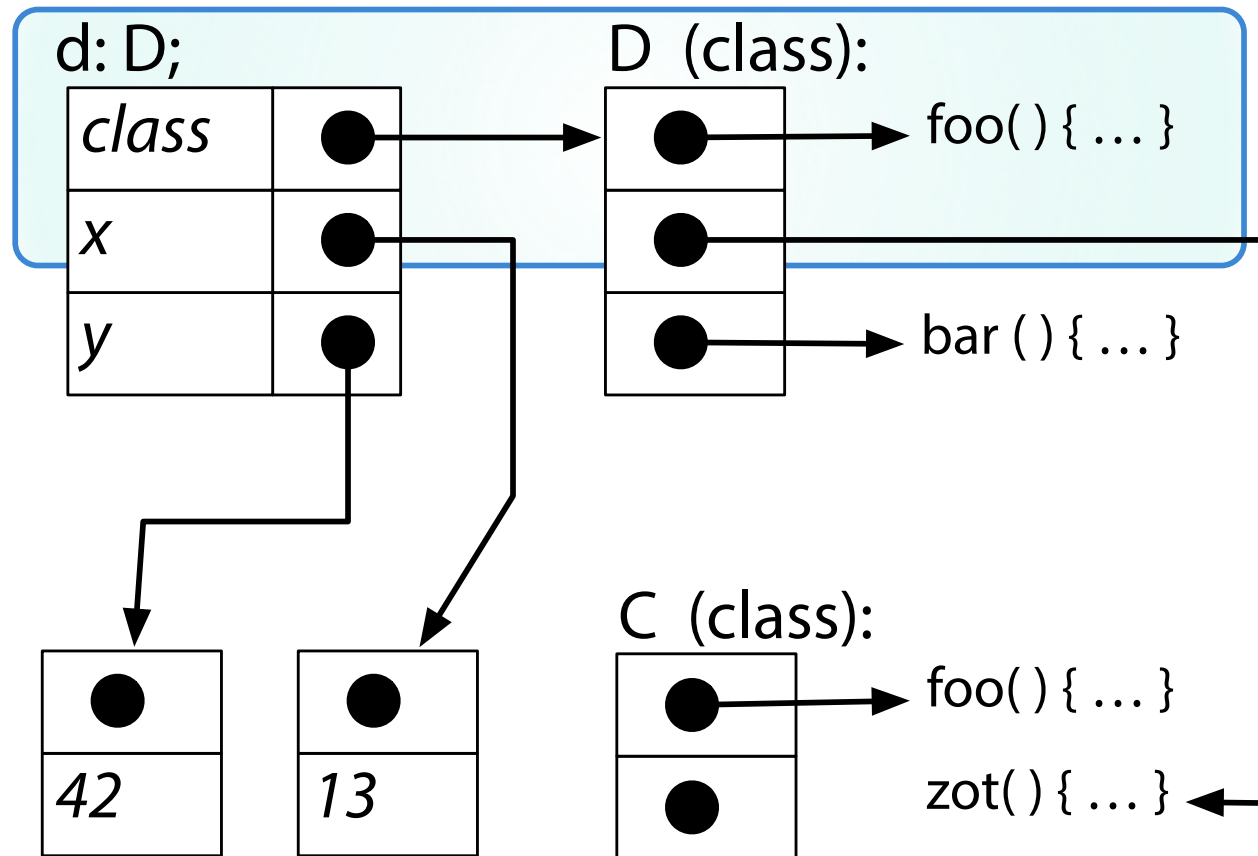


# Addressing fields (instance variables) in an object



Each field is at some fixed offset from the beginning of the object.  
If `D` is a subclass of `C`, its inherited fields are at the same offset.

# Method dispatch




Each method is at a known offset in the class object. We follow the object pointer, then index into the method table. Subclass method tables align with superclass method tables, so method offsets work for any subclass of static class.

In C: members of structs stand in for offsets in object. Subclasses have same member names at start of struct.



# Built-in classes

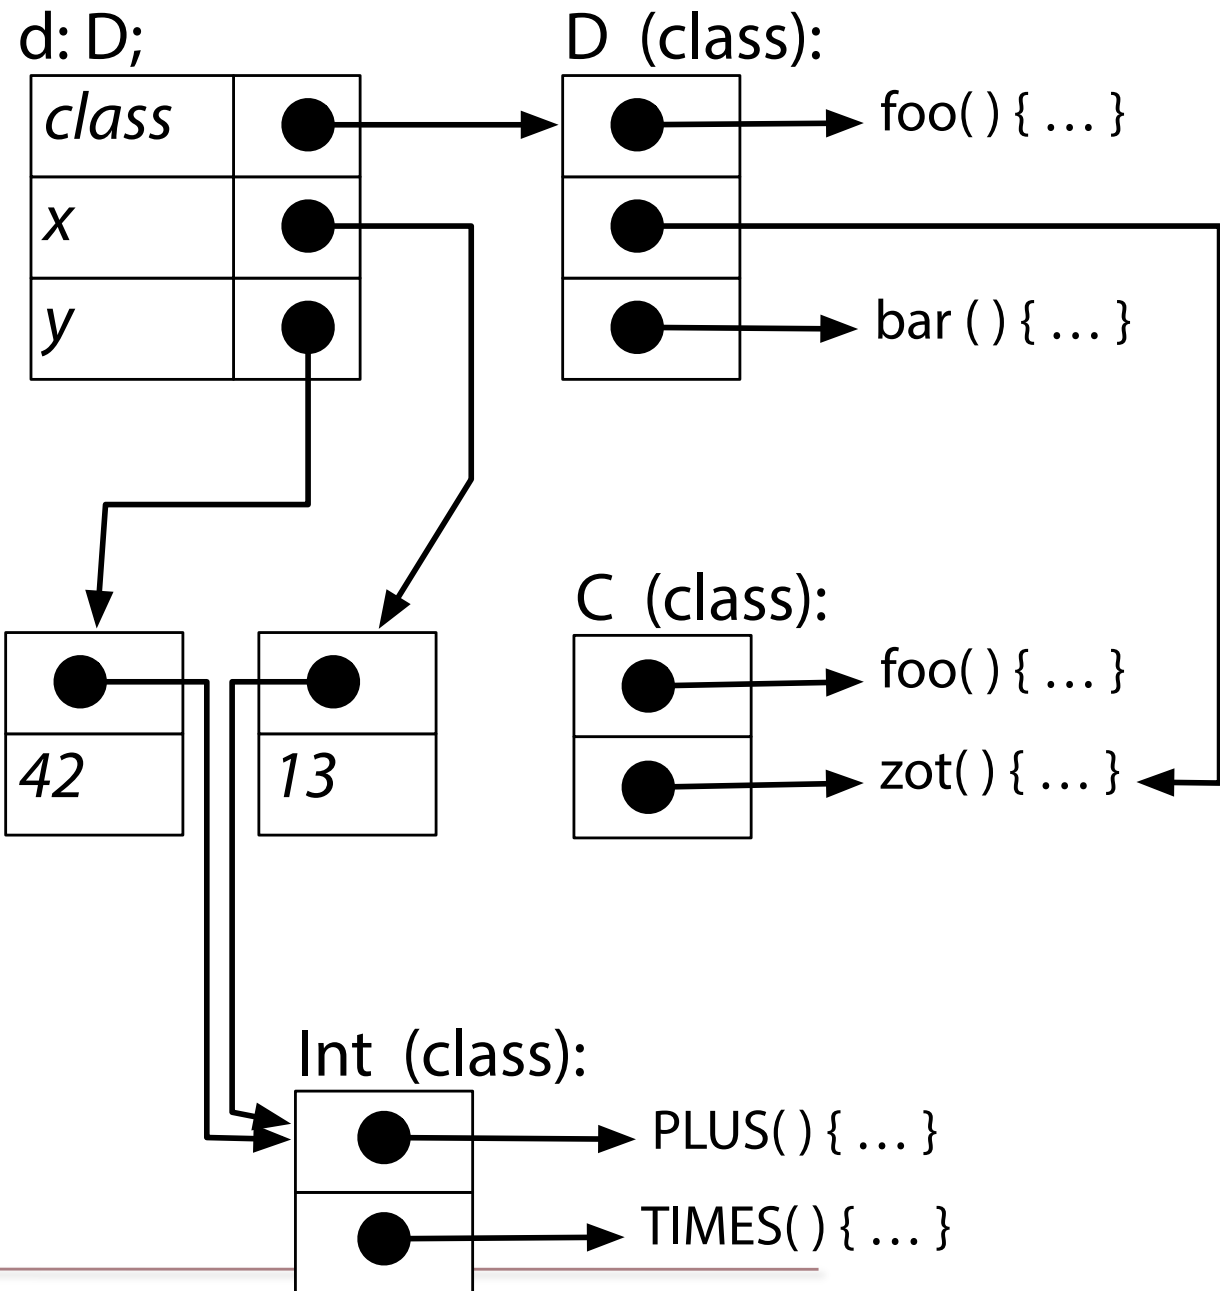
All values in Quack are objects. Even Int values are “boxed”. Only built-in object types may have non-object fields, which may be accessed by methods written in C. Primitive classes mimic Quack classes in layout, to minimize special cases.



## The whole picture, again ...

```
class C( ) {  
    this.x = 42;  
    def foo( ): Int {  
        this.zot( );  
    }  
    def zot( ) { ... }  
}
```

```
class D() extends C {  
    this.x = 43;  
    this.y = 13;  
    def foo( ): Int {  
        this.zot( );  
        this.bar( );  
    }  
    def bar( ) { ... }  
}
```



# *What About Expressions?*

Expression evaluation: Walk the tree

Leaves up, leaving result in register

Method calls: Evaluate each argument, place in temporary, get method address, make the call

Control structures - boilerplate

Conditionals

Short circuit evaluation

